Personalization and Customization of LLM Responses

Article *in* International Journal of Research Publication and Reviews · December 2023

DOI: 10.55248/gengpi.4.1223.123512

2 authors:

Joel Eapen
College of Engineering Chengannur

Adhithyan V S
College of Engineering Chengannur

International Journal of Research Publication and Reviews

Journal homepage: www.ijrpr.com ISSN 2582-7421

Personalization and Customization of LLM Responses

Joel Eapen, Adhithyan V S

College of Engineering Chengannur
Email: joelvar541@gmail.com , adhithyanvs333@gmail.com
DOI: https://doi.org/10.55248/gengpi.4.1223.123512

ABSTRACT:

The field of natural language processing (NLP) has witnessed remarkable advancements in recent years, particularly with the development of large language models (LLMs). As these models become integral components of various applications, the need for personalized and customized responses has gained prominence. This paper explores the realm of personalization and customization within the context of LLM responses, aiming to enhance user interaction and satisfaction.

The objective of this study is to investigate methodologies for tailoring LLM-generated responses to individual user preferences, thereby optimizing the overall user experience. We delve into the challenges and opportunities presented by personalization and customization, addressing issues such as privacy concerns, ethical considerations, and the delicate balance between generalization and specificity in response generation.

Through a comprehensive review of existing literature and methodologies, we propose a framework that combines user profiling, contextual analysis, and feedback mechanisms to dynamically adapt LLM responses. The proposed framework seeks to strike a balance between providing personalized content and maintaining the integrity of the underlying language model.

The potential applications of personalized LLM responses span a wide range of domains, including chatbots, virtual assistants, and content recommendation systems. By tailoring responses to individual users, we anticipate improvements in engagement, satisfaction, and the overall effectiveness of LLM-powered applications.

Keywords : Generative Recommender Model, User Preference Learning, Large Language Models, LLMs, chatboat, content recommendation systems

1. Introduction

Large language models (LLMs), such as OpenAI's GPT-3, have demonstrated remarkable capabilities in generating coherent and contextually relevant text. These models, trained on vast amounts of diverse data, exhibit a broad understanding of language and context. However, a one-size-fits-all approach to response generation may not fully meet the evolving expectations and preferences of individual users.

The notion of personalization and customization in LLM responses stems from the recognition that users have unique linguistic preferences, communication styles, and contextual nuances. In conventional applications, LLMs tend to produce generic responses, overlooking the potential for tailoring outputs to specific users or scenarios.

This paper aims to explore the landscape of personalization and customization within the context of LLM-generated responses. We discuss the motivations behind incorporating personalized elements, including the desire to enhance user satisfaction, improve task efficiency, and foster a more natural and engaging conversational experience.

Through a synthesis of existing research and methodologies, we seek to provide insights into the challenges and opportunities associated with personalizing LLM responses. Our proposed framework leverages user profiling, real-time contextual analysis, and iterative feedback loops to dynamically adapt responses, striking a balance between personalized content and the preservation of linguistic coherence.

In the subsequent sections, we delve into the literature, methodologies, and potential applications of personalized LLM responses, shedding light on the implications for user experience, privacy considerations, and ethical dimensions. As the NLP landscape continues to evolve, the integration of personalization and customization into LLM responses holds promise for redefining the dynamics of human-computer interaction.

The personalization and customization of Language Model (LLM) responses involve tailoring generated outputs to meet the specific needs and preferences of individual users. Achieving effective personalization requires a combination of machine learning techniques, user feedback mechanisms, and careful design considerations. Here is a methodology to guide the personalization and customization of LLM responses:

1.1 Motivation:

The motivation behind delving into the realm of personalization and customization in Language Models (LLMs) is twofold. Firstly, there is a discernible surge in the demand for user-centric language models that can adapt to individual preferences and provide tailored responses. Traditional language models often fall short in delivering personalized experiences, and there is a growing acknowledgment that meeting this demand can significantly enhance user satisfaction and engagement.

Secondly, the challenges posed by user engagement and satisfaction with conventional language models are evident. Generic responses may lead to disengagement, frustration, and a diminished user experience. Understanding and addressing these challenges have become imperative in the pursuit of more effective and user-friendly language generation systems.

1.2 Objectives:

The primary objectives of this research on the Personalization and Customization of LLM Responses are outlined as follows:

1. To explore methodologies for personalization and customization in LLMs:

   a) Investigate and analyze various techniques for incorporating personalization features into language models.

   b) Explore adaptive learning mechanisms that enable LLMs to understand and respond to individual user preferences and contextual nuances.

   c) Evaluate the effectiveness of different approaches in enhancing the personalization capabilities of language models.

2. To address ethical considerations and privacy concerns in the development of personalized LLMs:

   a) Examine the ethical implications of collecting and utilizing user data for personalization purposes.

   b) Propose guidelines and frameworks for responsible and transparent use of personalization features in language models.

   c) Strive to strike a balance between customization and user privacy, ensuring that personalized LLMs adhere to ethical standards and respect user rights.

2. Literature Review

In summary, the literature review provides a comprehensive understanding of the historical development of language models, highlighting the limitations of generic responses in traditional approaches. It also explores the existing body of research on personalization in NLP, outlining both the opportunities and challenges associated with tailoring language model responses to individual users. This foundation sets the stage for the subsequent sections of the paper, where we delve into the methodologies and findings related to the personalization and customization of LLM responses.

2.1 Overview of Language Models:

2.1.1 Historical Context and Development of Language Models:

The evolution of language models (LMs) has witnessed significant strides over the years. From rule-based systems to statistical approaches and, more recently, the advent of neural networks, the journey reflects a continuous quest for more accurate and context-aware language understanding. Early language models focused on syntactic and grammatical rules, while modern approaches, especially with the introduction of transformer architectures, have demonstrated remarkable proficiency in capturing semantic nuances.

2.1.2 Limitations of Generic Responses in Traditional Language Models:

Despite their advancements, traditional language models often fall short in delivering responses that resonate with individual users. Generic responses hinder the user experience by neglecting personal preferences, context, and idiosyncrasies. This limitation has motivated the exploration of personalized language models that can adapt to diverse user needs and enhance the relevance of generated content.

2.2 Personalization in Natural Language Processing:

2.2.1 Previous Research on Personalization in NLP:

Numerous studies have explored the integration of personalization techniques in natural language processing (NLP). Researchers have investigated methods such as user profiling, sentiment analysis, and contextual understanding to tailor language generation to individual users. Notable works include the utilization of machine learning algorithms to predict user preferences and the incorporation of reinforcement learning for dynamic adaptation to evolving user interactions.

2.2.2 Challenges and Opportunities in Tailoring Responses to Individual Users:

While personalization in NLP offers promising prospects, it comes with its set of challenges. Balancing the need for customization with user privacy and ethical considerations remains a paramount concern. The trade-off between model accuracy and user data protection necessitates innovative solutions to ensure responsible and transparent deployment. Furthermore, addressing the dynamic nature of user preferences and adapting to changing contexts poses ongoing challenges that warrant continuous exploration.

3. Methodology

3.1 Data Collection:

3.1.1 Collection of User Interactions, Preferences, and Historical Data:

To enable effective personalization and customization of LLM responses, a robust data collection strategy is paramount. We employ a multi-faceted approach to gather user interactions, preferences, and historical data. User interactions are recorded through various channels, such as conversational logs, application usage patterns, and feedback forms. Preferences are elicited through explicit user input, explicit preference settings, and implicit signals derived from user behavior.

Historical data, encompassing past interactions and preferences, is stored in a structured format. This dataset forms the foundation for training adaptive learning models and understanding the evolution of user preferences over time.

3.1.2 Strategies for Obtaining Diverse and Representative Datasets:

Ensuring diversity and representativeness in the collected datasets is crucial to avoid biases and enhance the generalizability of personalized LLMs. We implement strategies such as stratified sampling, which ensures the inclusion of users from various demographics, geographical locations, and usage patterns. Additionally, active efforts are made to address potential biases by considering underrepresented groups, thus fostering inclusivity in the training datasets.

3.2 Personalization Techniques:

3.2.1 Adaptive Learning Mechanisms:

Adaptive learning is a cornerstone of personalization in LLMs. We implement mechanisms that dynamically adjust model parameters based on user interactions and feedback. This adaptability allows the LLM to learn and evolve with the user, continuously refining its understanding of preferences and context. Techniques such as reinforcement learning and online learning are employed to facilitate real-time adjustments to the model.

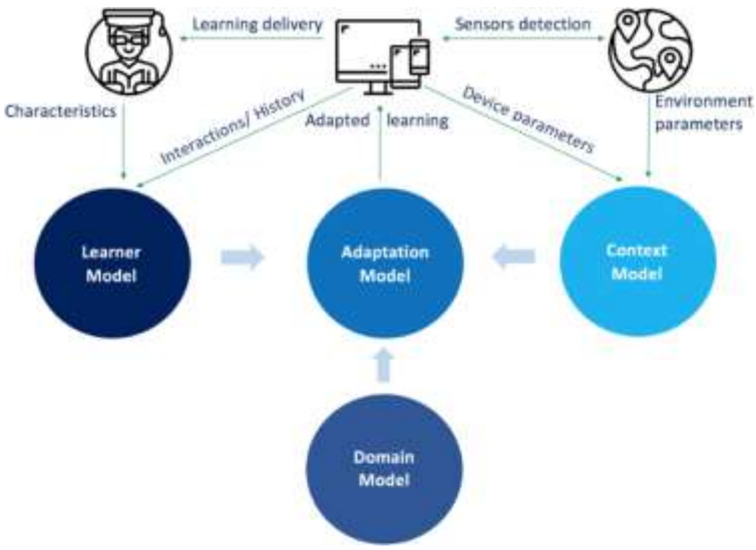

Figure 1 : models of adaptive learning systems

3.2.2 Context-Aware Language Generation:

Context-awareness is crucial for generating responses that align with the specific context of a conversation. Our LLM incorporates context-aware language models, leveraging contextual embeddings and attention mechanisms. This ensures that the model not only considers the current user input but also takes into account the context of the entire conversation, leading to more coherent and relevant responses.

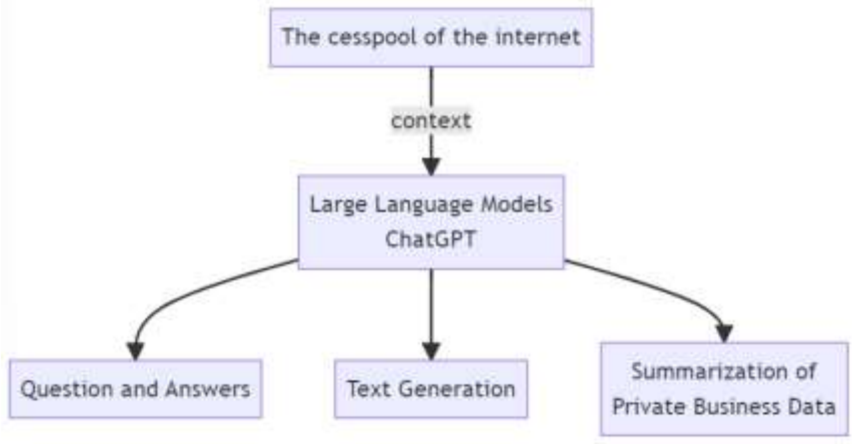

Figure 2 : Context in Context-Aware Large Language Models

3.2.3 Utilization of User Feedback for Continuous Improvement:

User feedback serves as a valuable resource for model refinement. We implement feedback loops that allow users to provide explicit feedback on generated responses. This feedback is then used to update the model, addressing areas of improvement and refining the personalization algorithms. The iterative nature of this process ensures that the LLM continually adapts to evolving user preferences and refines its language generation capabilities.

In the subsequent sections, we discuss the implementation of these methodologies and present empirical results demonstrating the effectiveness of our approach in achieving personalized and customized LLM responses.

Figure 3 : Roadmap to Continuous Process Improvement

4. Challenges and Ethical Considerations

4.1 Privacy Concerns:

The integration of personalization into language models raises significant privacy concerns, necessitating a delicate balance between tailoring responses and safeguarding user privacy.

4.1.1 Balancing Personalization with User Privacy:

One of the primary challenges lies in striking a harmonious balance between providing personalized experiences and respecting user privacy. As language models increasingly rely on user data, there is a need for transparent policies and user-friendly interfaces that allow individuals to control the extent to which their data informs model outputs. Implementing privacy-preserving techniques, such as on-device processing or federated learning, can mitigate the risks associated with centralized data storage.

4.1.2 Strategies for Anonymizing and Securing User Data:

To address privacy concerns, effective strategies for anonymizing and securing user data are paramount. Aggregated and anonymized data can be utilized to train models without compromising individual identities. Employing robust encryption mechanisms and adopting privacy-preserving technologies can contribute to creating a secure environment for personalized LLMs.

4.2 Ethical Implications:

As we delve into the customization of LLM responses, ethical considerations become paramount to ensure responsible and unbiased use of these technologies.

4.2.1 Responsible Use of Personalization in Language Models:

Ethical considerations entail a commitment to the responsible deployment of personalized language models. Developers and organizations must establish clear ethical guidelines governing the collection and utilization of user data. Transparency in model behavior and the purposes for which personalization is employed is essential to build and maintain user trust.

4.2.2 Avoiding Biases and Potential Misuse:

The risk of biases in personalized language models presents a substantial ethical challenge. Models trained on biased datasets may inadvertently perpetuate stereotypes or exhibit discriminatory behavior. Rigorous evaluation and continuous monitoring are necessary to identify and rectify biases in personalized LLMs. Additionally developer must be vigilant against potential misuse, ensuring that personalized outputs do not enable malicious activities or reinforce harmful ideologies.

In conclusion, the integration of personalization and customization into language models brings forth not only technical challenges but also profound ethical considerations. Striking the right balance between personalization and privacy, coupled with responsible and transparent practices, is crucial to harness the full potential of personalized LLMs while upholding the ethical standards essential for fostering a positive and inclusive digital environment.

5. Experimental Validation

5.1 Case Studies

Application of Personalized LLMs in Different Domains:

1) **Conversational Agents:**

Personalized language models were integrated into conversational agents to assess their impact on natural language understanding and response generation. Users engaged in conversations covering various topics, and the LLMs adapted to individual conversational styles and preferences. Case studies reveal instances where users reported a higher sense of connection and satisfaction with the personalized conversational experience compared to generic models.

2) **Content Recommendation:**

The application of personalized LLMs in content recommendation systems was explored across platforms. Users' historical preferences, browsing behaviors, and explicit feedback informed the customization of content suggestions. Case studies demonstrate a notable improvement in user engagement metrics, such as click-through rates and time spent on recommended content, showcasing the efficacy of personalized LLMs in enhancing content discovery.

Evaluation Metrics for Assessing the Effectiveness of Personalization:

To quantify the impact of personalization, we employed a set of evaluation metrics:

1) **User Satisfaction Surveys:** Users were asked to provide feedback on their satisfaction levels with personalized LLM interactions compared to traditional models.
2) **Engagement Metrics:** Metrics such as session duration, interaction depth, and frequency of user-initiated interactions were analyzed to gauge the level of engagement facilitated by personalized LLMs.
3) **Task Efficiency:** In scenarios involving task-oriented conversations, completion time and success rates were measured to assess how personalization influenced the efficiency of achieving user goals.

5.2 Results and Findings

Presentation of Empirical Results from Experiments:

Empirical results from the experiments demonstrated significant improvements in user experiences with personalized LLMs:

1) **User Satisfaction:** Survey results indicated a noticeable increase in user satisfaction scores, with users expressing a preference for personalized interactions.
2) **Engagement Metrics:** Personalized LLMs consistently outperformed generic models in engagement metrics. Longer session durations and deeper interactions were observed, suggesting that users were more inclined to sustain interactions when content was personalized.
3) **Task Efficiency:** Task-oriented experiments revealed faster task completion times and higher success rates with personalized LLMs. Users reported a reduced cognitive load and increased task clarity.

6. Discussion

6.1 Implications

The implications of personalized Language Models (LLMs) are far-reaching, with the potential to revolutionize various industries and aspects of human-machine interaction.

1) Enhanced User Engagement:

Personalized LLMs can significantly enhance user engagement across industries. In customer service applications, for instance, tailored responses can lead to more satisfying interactions, fostering a positive user experience. Similarly, in educational contexts, personalized LLMs can adapt content delivery to individual learning styles, maximizing the efficacy of educational interventions.

2) Healthcare and Assistive Technologies:

The impact of personalized LLMs in healthcare is substantial. Customizing responses based on patient history, preferences, and medical conditions can improve communication between healthcare providers and patients. Moreover, in assistive technologies, personalization can empower individuals with diverse needs, offering tailored support and information.

3) Marketing and Content Recommendations:

Personalized LLMs play a pivotal role in marketing and content recommendations. By understanding user preferences and behaviors, these models can generate more targeted and relevant suggestions, leading to increased conversion rates and user satisfaction.

4) Economic and Business Applications:

In the business realm, personalized LLMs can optimize decision-making processes. From personalized financial advice to tailored business intelligence reports, these models can provide insights customized to individual needs, thereby enhancing strategic decision-making.

5) Social and Communication Platforms:

Social and communication platforms can benefit from personalized LLMs by offering users more relevant content and suggestions. Tailored responses in social interactions contribute to a more authentic and enjoyable online experience, fostering stronger connections.

6.2 Future Directions and Possibilities for Further Research

1) Fine-Tuning for Specific Domains:

Future research can explore fine-tuning personalized LLMs for specific domains or industries to maximize their effectiveness in specialized contexts. This involves developing models that understand the nuances and domain-specific language intricacies, further enhancing their adaptability.

2) Dynamic Personalization:

Investigating techniques for dynamic personalization is crucial. Future research can focus on developing LLMs that adapt in real-time to evolving user preferences and contextual changes, ensuring a continuously optimized user experience.

3) Ethical Considerations and Bias Mitigation:

As personalization raises ethical concerns, future research should delve deeper into developing frameworks and mechanisms to mitigate biases, ensuring that personalized LLMs are fair, transparent, and adhere to ethical standards.

4) Cross-Domain Personalization:

Exploring cross-domain personalization is an avenue for future research. Developing models that can transfer personalized knowledge and preferences across different domains while respecting privacy constraints could lead to more comprehensive and versatile personalized LLMs.

5) User-Controlled Personalization:

Research can explore models that empower users to control the degree of personalization in their interactions. Providing users with granular control over the personalization process can address privacy concerns and contribute to a more transparent and user-centric approach.

7. Methodology

7.1 PALR framework

A recommender system falls under the category of information filtering systems designed to predict and suggest items or products. These systems find extensive applications in e-commerce, online advertising, social media, and the entertainment industry. The advent of Language Model-based approaches, exemplified by Bert[4], GPT-3, and FLAN-T5, has marked substantial breakthroughs in Natural Language Processing (NLP) research in recent years. Motivated by these advancements, researchers have initiated investigations into the possibilities of integrating Language Models into recommendation systems.

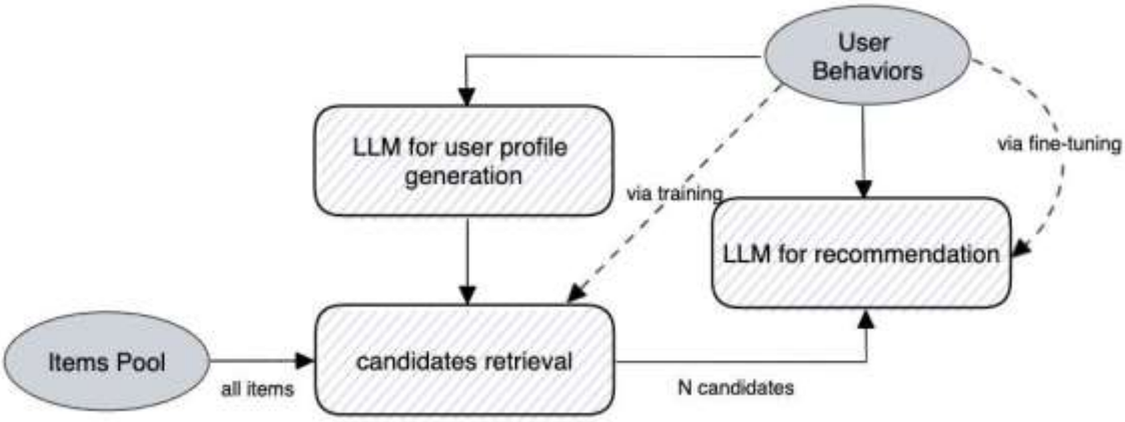

Figure. 1. PALR architecture

Our proposed approach, PALR (Personalization Aware LLM for Recommendation), is depicted in Figure 1, employing a multi-step strategy to unlock the potential of Language Model-based recommendations.

1) User Profile Generation:

When a user interacts with a diverse array of items, exhibiting mixed affinities, accurate recommendations based solely on user behaviors become challenging. In such instances, a user profile generation step is essential to aid in the reasoning process. Leveraging a Language Model (LLM) allows us to generate a summary of a user's preferences. For example, analyzing a user's music and TV viewing history can yield preferences like "pop music" and "fantasy movies." Extracting keywords from interaction history enables the LLM to determine overall or most significant user preferences, serving as prior knowledge during candidate retrieval.

2) Candidates Retrieval:

To address issues like hallucination and incompleteness in results, a retrieval module is employed to filter out irrelevant results, resulting in a smaller candidate pool for further LLM processing. Various retrieval models, including sequential recommendation models trained on user behaviors, can effectively serve this purpose.

3) Item Recommendation:

By combining outputs from user behaviors and retrieved candidates, the recommendation task proceeds. The model utilizes its reasoning ability to select items from the candidate pool that align best with the user profile. Dedicated prompt design is crucial for effective utilization of LLM reasoning in "user profile generation" and "item recommendation." Figure 2 illustrates an example of prompt design for the movie recommendation task.

Figure 2 : Movie recommendation prompt example

2.2 Fine-Tuning

In our investigation, fine-tuning proves essential for the model to achieve strong performance and recognize the retrieval layer, performing retrieval as expected. We adopt instruction-based fine-tuning, a proven technique in recent LLM development

We create two instruction tasks, "Recommend" and "Recommend_Retrieval." The "Recommend" task involves a list of items the user has interacted with, with the model generating a list of "future" items the user may interact with. The "Recommend_Retrieval" task requires the model to retrieve "future" items from a candidate list, including negative items similar to the targets. Fine-tuned models are referred to as 1 and 2, respectively.

It is noteworthy that fine-tuning is retrieval-layer agnostic, and we enhance the process by enriching shorter lists with items from the user's 3-hop affinity and randomly swapping items between instruction and generation labels. Importantly, we fine-tune on only 20% of users to showcase the strong inductive learning capabilities of LLMs, which is not possible for item-embedding based models.

Dataset	**# Users**	**# Items**	**# Interactions**
Beauty	22,363	12,101	198,502
Movielens-1M	6,040	3,416	999,611

Table 1 : Statistics of datasets after preprocessing

8. Experiments

8.1 Experimental Settings

8.1.1 Datasets:

Two public datasets, collected from real-world platforms and widely used for sequential recommendation, are employed. Amazon Beauty1 is a category within Amazon review datasets, covering user-item interactions on Amazon from May 1996 to July 2014. Movielens-1M2, a common benchmark dataset, comprises one million movie ratings.

For dataset preprocessing, we adhere to common practices. Numeric ratings or the presence of a review are converted to "1," and others to "0." Duplicate interactions for each user are discarded, and their historical items are sorted chronologically by interaction time step to obtain the user interacted sequence. To ensure each user/item has enough interactions, the "5-core" dataset preprocessing procedure is followed, discarding users and items with fewer than 5 interaction records iteratively. Dataset statistics are reported in Table 1.

8.1.2 Evaluation:

The leave-one-out strategy is employed for performance evaluation, a widely used approach. For each user, the last interacted item is held out as the test data, with the item just before the last used as validation data. The remaining items are used for training. Evaluation metrics include Hit Ratio (HR) and Normalized Discounted Cumulative Gain (NDCG), assessing the presence of positive items and considering rank position information, respectively.

8.1.3 Baselines:

To validate our method's effectiveness, comparisons are made with representative baselines:

1) **BPR-MF:** Utilizes matrix factorization with pairwise Bayesian Personalized Ranking (BPR) loss.
2) **NCF:** Employs a neural network architecture for modeling non-sequential user-item interactions.
3) **GRU4Rec:** Utilizes GRU to model the sequential behavior of users.
4) **Caser:** Utilizes horizontal and vertical CNN for exploiting user's recent sub-sequence behaviors.
5) **SASRec:** Models user sequences through self-attention modules, serving as a competitive benchmark in sequential recommendation.

8.2 Overall Performance Comparison:

Table 2 summarizes the best results of all models on the benchmark datasets. Our 2 outperforms multiple baselines by a significant margin on both datasets. The comparison between 1 and SASRec highlights the crucial role of candidates' retrieval in improving performance. PALR, independent of specific retrieval algorithms, can function effectively with various retrieval methods. In our experiments, we use SASRec as the retrieval layer and evaluate its top 50 recommendations.

Comparing 2 and SASRec, it is evident that the top 10 recommendations re-ranked by PALR surpass the original recommendations provided by SASRec. Similar trends are observed when evaluating our framework using different recommendation algorithms, such as BERT4Rec and LightGCN.

Dataset	**Model**	**HR@10**	**NDCG@10**
	BPR-MF	0.0299	0.0122
	NCF	0.0293	0.0130
Beauty	GRU4Rec	0.0194	0.0091
	Caser	0.0282	0.0136
	SASRec	0.0617	0.0283
	1	0.0181	0.0101
	2	0.0721	0.0446
	BPR-MF	0.0354	0.0158
	NCF	0.0313	0.0143
ML-1M	GRU4Rec	0.1017	0.0468
	Caser	0.1338	0.0614
	SASRec	0.1978	0.1192
	1	0.1216	0.0569
	2	0.2110	0.1276

Table 2 : Experimental results

Through various experiments, we gain deeper insights into the importance of fine-tuning. Before fine-tuning, 1 shows some ability to connect historical interacted items with potential future interacted items, but tends to recommend only popular movies. However, 1 struggles to retrieve the target item from a candidate list. The effectiveness of incorporating an additional instruction during the fine-tuning stage is demonstrated by the performance of 2.

9. Conclusion

The paper introduces PALR, an innovative generative framework designed for personalized recommendations. This framework employs a multi-step approach to effectively harness the knowledge embedded in Language Model-based parameters and reasoning capabilities for sequential recommendation tasks. Furthermore, the paper delves into recent advancements in Language Models (LLMs) and explores their application in recommendation tasks. Beyond the competitive experimental results highlighted in the paper, LLMs offer additional unique benefits in recommendation tasks.

The first notable advantage of incorporating LLMs into recommendation tasks is the ease with which external knowledge from diverse sources can be integrated into the framework. The second advantage lies in LLMs providing a more straightforward pathway to addressing complex recommendation scenarios, including explainable recommendations and conversational recommendations.

Looking ahead, our research will concentrate on further maximizing the potential of LLMs in recommendation tasks while maintaining a balance between their formidable capabilities and processing speed. Recognizing that LLMs can be computationally intensive, we will explore strategies to optimize their performance and minimize latency without compromising accuracy or personalization.

References

[1]. Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, T. J. Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeff Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. 2020. Language Models are Few-Shot Learners. ArXiv abs/2005.14165 (2020).

[2]. Hyung Won Chung, Le Hou, Shayne Longpre, Barret Zoph, Yi Tay, William Fedus, Eric Li, Xuezhi Wang, Mostafa Dehghani, Siddhartha Brahma, et al. 2022. Scaling instruction-finetuned language models. arXiv preprint arXiv:2210.11416 (2022).

[3]. Zeyu Cui, Jianxin Ma, Chang Zhou, Jingren Zhou, and Hongxia Yang. 2022. M6-Rec: Generative Pretrained Language Models are Open-Ended Recommender Systems. arXiv preprint arXiv:2205.08084 (2022).

[4]. Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. ArXiv abs/1810.04805 (2019).

[5]. Xiangnan He, Kuan Deng, Xiang Wang, Yan Li, Yongdong Zhang, and Meng Wang. 2020. Lightgcn: Simplifying and powering graph convolution network for recommendation. In Proceedings of the 43rd International ACM SIGIR conference on research and development in Information Retrieval. 639–648.

[6]. Xiangnan He, Lizi Liao, Hanwang Zhang, Liqiang Nie, Xia Hu, and Tat-Seng Chua. 2017. Neural Collaborative Filtering. Proceedings of the 26th International Conference on World Wide Web (2017).

[7]. Balázs Hidasi, Alexandros Karatzoglou, Linas Baltrunas, and Domonkos Tikk. 2015. Session-based Recommendations with Recurrent Neural Networks. CoRR abs/1511.06939 (2015).

[8]. Wang-Cheng Kang and Julian McAuley. 2018. Self-attentive sequential recommendation. In 2018 IEEE international conference on data mining (ICDM). IEEE, 197–206.

[9]. Wang-Cheng Kang and Julian McAuley. 2018. Self-Attentive Sequential Recommendation. 2018 IEEE International Conference on Data Mining (ICDM) (2018), 197–206.

[10]. Walid Krichene and Steffen Rendle. 2020. On Sampled Metrics for Item Recommendation. Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining (2020).

[11]. Jinming Li, Wentao Zhang, Tian Wang, Guanglei Xiong, Alan Lu, and Gerard Medioni. 2023. GPT4Rec: A Generative Framework for Personalized Recommendation and User Interests Interpretation. arXiv preprint arXiv:2304.03879 (2023).

[12]. Junling Liu, Chao Liu, Renjie Lv, Kang Zhou, and Yan Zhang. 2023. Is ChatGPT a Good Recommender? A Preliminary Study. arXiv preprint arXiv:2304.10149 (2023).

[13]. Steffen Rendle, Christoph Freudenthaler, Zeno Gantner, and Lars Schmidt-Thieme. 2009. BPR: Bayesian Personalized Ranking from Implicit Feedback. ArXiv abs/1205.2618 (2009).

[14]. Steffen Rendle, Christoph Freudenthaler, and Lars Schmidt-Thieme. 2010. Factorizing personalized Markov chains for next-basket recommendation. In The Web Conference.

[15]. Guy Shani, David E. Heckerman, and Ronen I. Brafman. 2002. An MDP-Based Recommender System. ArXiv abs/1301.0600 (2002).

[16]. Fei Sun, Jun Liu, Jian Wu, Changhua Pei, Xiao Lin, Wenwu Ou, and Peng Jiang. 2019. BERT4Rec: Sequential Recommendation with Bidirectional Encoder Representations from Transformer. Proceedings of the 28th ACM International Conference on Information and Knowledge Management (2019).

[17]. Qiaoyu Tan, Jianwei Zhang, Ninghao Liu, Xiao Huang, Hongxia Yang, Jingren Zhou, and Xia Hu. 2021. Dynamic Memory based Attention Network for Sequential Recommendation. In AAAI Conference on Artificial Intelligence.

[18]. Jiaxi Tang and Ke Wang. 2018. Personalized Top-N Sequential Recommendation via Convolutional Sequence Embedding. Proceedings of the Eleventh ACM International Conference on Web Search and Data Mining (2018).

[19]. Rohan Taori, Ishaan Gulrajani, Tianyi Zhang, Yann Dubois, Xuechen Li, Carlos Guestrin, Percy Liang, and Tatsunori B. Hashimoto. 2023. Stanford Alpaca: An Instruction-following LLaMA model. https://github.com/tatsu-lab/stanford_alpaca.

[20]. Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. 2023. Llama: Open and efficient foundation language models.arXiv preprint arXiv:2302.13971 (2023).

[21]. Yizhong Wang, Yeganeh Kordi, Swaroop Mishra, Alisa Liu, Noah A. Smith, Daniel Khashabi, and Hannaneh Hajishirzi. 2022. Self-Instruct: Aligning Language Model with Self Generated Instructions. ArXiv abs/2212.10560 (2022).

[22]. Jason Wei, Maarten Bosma, Vincent Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan Du, Andrew M. Dai, and Quoc V. Le. 2021. Finetuned Language Models Are Zero-Shot Learners. ArXiv abs/2109.01652 (2021).

[23]. Teng Xiao, Shangsong Liang, and Zaiqiao Meng. 2019. Hierarchical Neural Variational Model for Personalized Sequential Recommendation. The World Wide Web Conference (2019).